Glossary

Terms whose meaning holds across the reports. A report shows each one beside its first use in every section, and a report that uses a term in a narrower sense defines it again in its own glossary.

Alignment

α

How far a state points along e_1: the cosine between the state and e_1, which for a unit-length state is its component along e_1; 1 on the anchor direction and 0 off it.

Bayes predictor

The best answer given the examples: it weighs the answers of each op by the posterior over ops, the probability of each op given the examples.

Expected exact match

The task score, as the chance that a drawn answer agrees with a true one. With drawn answers, the correct answer to a line is spread over two or more colors, so we take the chance that an answer drawn from the model agrees with one drawn from the true distribution. Its ceiling is that same chance for the true distribution against itself (Σq^2), below 1 on the ops that round.

Band

Our measurement precision, the smallest gap between seed means that we call resolved. Two seed means are told apart only when they differ by more than $2\sigma\sqrt{(1/n_a + 1/n_b)}$, with σ the per-run spread and n the seed counts; smaller differences are reported as unresolved.

Pooling temperature

How sharply the anchor pull picks its positions. The pull on a context is a softmin-weighted mix over its positions, and the temperature (τ_p) sets how peaked the weights are: at 0.1 nearly all the weight goes to the position closest to e_1, and at 1 it is spread over many.

Spill

The side effect of an edit on the other ops: the largest drop in the score of any op other than the one removed, net of what the same edit does to a control trained without the anchor. The selectivity criterion of D2.2 allows at most 0.02 at any dose.

Target null

Where a removal should land: what an ideal predictor that had lost one op and nothing else would answer. It is the Bayes posterior over ops given the examples, with the removed op set to zero and the rest renormalized, so it still weighs the other ops by how well they fit.

Total variation

Total variation distance

A distance between two answer distributions: half the summed absolute difference, from 0 (identical) to 1 (no overlap). Mass a model puts on tokens that are not colors counts as distance.